

PRIYA VENKATARAMAN

Senior Data Engineer

priya.venkat@gmail.com | (408) 555-0136 | San Jose, CA

<https://linkedin.com/in/priya-venkataraman> | <https://github.com/pvenkat>

Portfolio: www.priyavenkat.dev

SUMMARY

Data engineer with nine years building streaming and batch platforms on AWS and GCP. Led teams of 3-6; deep experience with Kafka, Flink, dbt, and lakehouse table formats. Comfortable owning systems from schema design through on-call.

EXPERIENCE

Senior Data Engineer · Lumen Analytics

500 Almaden Blvd, San Jose, CA 95110 | Mar 2021 - Present

- Built streaming ingestion handling 2B events/day on Kafka and Flink with exactly-once semantics into Iceberg.
- Led migration of nightly batch ETL to incremental pipelines, cutting warehouse spend 38% and P95 landing latency from 7h to 40min.
- Designed the client-facing usage-metering pipeline that bills \$40M ARR; zero billing incidents across eight quarters.
- Mentored four engineers; ran the data platform on-call rotation.

Data Engineer · Halcyon Metrics

221 River Oaks Pkwy, San Jose, CA 95134 | Jun 2018 - Feb 2021

- Owned the warehouse dimensional model (Snowflake, dbt) covering product, growth, and finance domains.
- Introduced data contracts between services and the warehouse; schema-drift incidents fell from ~3/month to <1/quarter.
- Built CDC ingestion from Postgres with Debezium.

Analytics Engineer · Corvid Software

Remote (Austin, TX HQ) | Jul 2016 - May 2018

- First data hire: stood up the warehouse, BI layer, and metrics definitions used company-wide.

SELECTED PROJECTS

-
- flink-iceberg-sink (maintainer): exactly-once Iceberg sink used by ~200 orgs. <https://github.com/pvenkat/flink-iceberg-sink>
 - Talk: "Metering at 2B events/day", Data Council 2024 · slides at <https://www.priyavenkat.dev/talks/data-council-2024>

EDUCATION & CERTIFICATIONS

B.S. Computer Science, San Jose State University, 2016

AWS Certified Data Engineer - Associate, issued 2025-01-10

Google Professional Data Engineer, issued 15-Mar-2023

SKILLS

Languages: Python, SQL, Java, Scala. Streaming: Kafka, Flink, Kinesis. Storage & modeling: Iceberg, Snowflake, BigQuery, dbt, Postgres. Orchestration: Airflow, Dagster. Infra: Terraform, Kubernetes, GitHub Actions. Observability: Monte Carlo, OpenLineage, Grafana.

TEACHING & COMMUNITY

-
- Instructor, "Practical Stream Processing", San Jose State continuing education, spring 2024 and spring 2025.
 - Co-organizer, South Bay Data Engineering meetup (~800 members); CFP and speaker liaison since 2022.
 - Reviewer, Journal of Open Source Software (data infrastructure track).

REFERENCES

Dr. Alan Reyes · Director of Engineering, Lumen Analytics. alan.reyes@lumenanalytics.com | +1 408-555-0171
Mei-Ling Zhou · VP Data, Halcyon Metrics. mzhou@halcyonmetrics.com | (669) 555-0148